

The WTCHG Research Computing Core

Talk 3: Submitting jobs to the cluster

Robert Esnouf (robert@strubi.ox.ac.uk; robert@well.ox.ac.uk),
Jon Diprose (jon@well.ox.ac.uk) & Colin Freeman (cfreeman@well.ox.ac.uk)

Generic emails: support rescomp@well.ox.ac.uk
users rescomp-users@well.ox.ac.uk

A series of introductory talks...

A set of six talks of about 45 minutes each (plus questions)

• Talk 1: What is the ResComp Core?

(Mon 23/1 10:00 Room B; Robert Esnouf)

Talk 2: A basic introduction to Linux

(Wed 25/1 10:00 Room B; Robert Esnouf)

Talk 3: Submitting jobs to the cluster

(Thu 26/1 10:00 Room B; Robert Esnouf)

Talk 4: Monitoring and troubleshooting

(Mon 30/1 11:00 Room B; Robert Esnouf)

Talk 5: ResComp centrally-managed applications

(Wed 1/2 10:00 Room B; Jon Diprose)

Talk 6: Doing your own thing (compiling and customizing)

(Thu 2/2 10:00 Room B; Jon Diprose)

The story so far...

Cluster Computing

- A nodes: 288 cores @ 2.67GB/core (32GB/node; 70% C speed)
- B nodes: 640 cores @ 8GB/core (96GB/node; 70% C speed)
- C nodes: 1720 cores @ 16GB/core (256GB/node; 100% C speed)
- D nodes: 768 cores @ 16GB/core (384GB/node; 100% C speed)
- H nodes: 96 cores @ 42.66GB/core (2TB/node; 115% C speed)

Login/submission nodes: rescomp[1-2] and the project servers

- General use: compiling programs, submitting jobs, analysing results

High-performance storage (GPFS)

- /gpfs0: 6.7 TB @ 400MB/s: */users/<group>/<user>, /apps/well, /mgmt*
- /gpfs1 & /gpfs2: 3075 TB @ >20GB/s: */well/<group>*
- Charging rate £35 / usable TB / 6 months

“Archive” Storage (XFS)

- /arc[1-3][a-g]: 1417TB read-only access. £14.16 / usable TB / 6 months

The story so far...

Basic refresher on key aspects of Linux

- User accounts, groups, UIDs and GIDs
- Commands, shells (bash) and environments
- Environment variables: \$PATH and \$LD_LIBRARY_PATH
- Standard file streams: stdin, stdout and stderr
- Redirection (“<”, “>”) and pipes (“|”)
- Foreground and background jobs (“&”), “nohup” and “screen”

Getting ready for using the cluster

- Writing a shell script
- Executing and sourcing a script
- Keeping track of stdin, stdout and stderr
- Cluster computing is essentially non-graphical and non-interactive!

Overview of this talk...

- Sun Grid Engine (SGE)

- Cluster queues and limits

- An SGE user

- My first cluster job: using qsub

- qsub options

- A first look at qstat

- Priority and the share tree

- SGE monitoring tools

Sun Grid Engine (SGE)

Sun Grid Engine (v. 6.2u5p3) schedules jobs across the cluster

- Free-to-use scheduler (and very simple resource manager)
- Familiar and simple: CLUSTER1, CLUSTER2 & CLUSTER3
- Not developed since Oracle bought Sun (2010) so we are stuck with bugs/features
- Alternatives LSF, PBS, PBSpro, Torque/Maui, slurm...
- “Son of Grid Engine” and Open Grid Engine (Kwiatkowski)
- Univa Grid Engine (UGE) is the commercial fork of SGE about £8-10k/yr for a site licence

Sun Grid Engine concepts

Conceptually very simple and resilient to failures

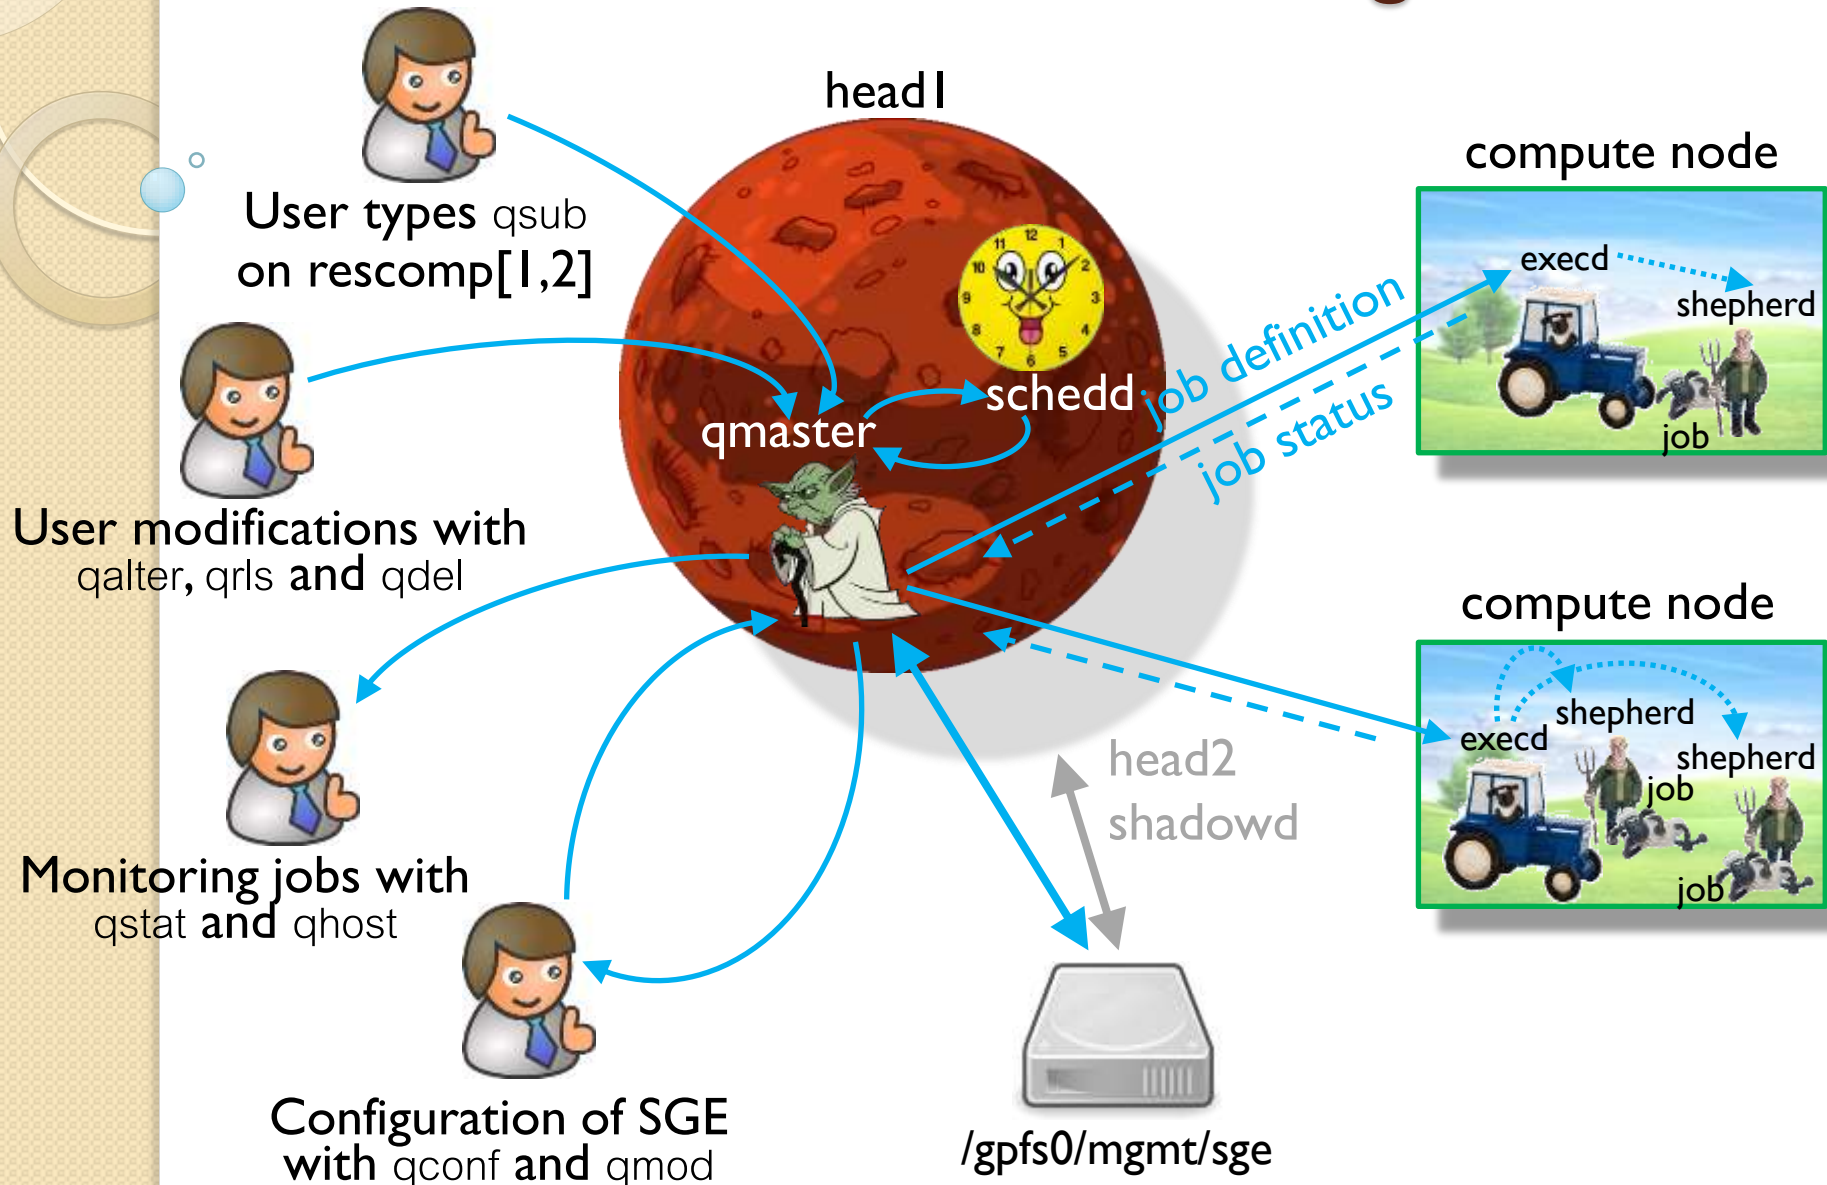
Processes:

- qmaster (on one head node)
- shadowd can start backup qmaster (on other head node)
- qmaster accepts job requests and monitoring requests
- qmaster fires off schedd for scheduling
- qmaster hands highest-priority job to execd on a compute node
- execd creates a job shepherd, detached from execd to run job

No state is held by these processes

- Always written to and/or read from filesystem (/gpfs0/mgmt/sge)
- /gpfs0 is a parallel file system mounted on both head1 and head2
- qmaster, schedd and execd do not need to be running to keep jobs running

Process flow with Sun Grid Engine



Cluster queues and limits

A queue is a definition of a set of rules for running a job

- Name of queue, list of hosts, who has access
- Number of slots (usually 1 per core)
- Run time limits, memory limits etc.

Mostly we define two queues per node type

- Short jobs (<24h) and long jobs (<7d)
- 1 core per slot, no use of resource limits (“-l”)

Special queues

- For relion jobs, for high memory jobs and on project servers

A queue instance is a queue on a single compute host

- The queue definition defines which queues are on hosts or host groups
- The `execd` on each host instantiates the queues on that host

General queues on the ResComp cluster

We can (& do) put both short & long queues on a single node

- A resource quota set (rqs) rule limits the total number of queue slots for all queues to the total number of cores on a node

Normal work queues

- short.qa and long.qa: 24h & 7d jobs on “A” nodes, 2.5 GB per slot
- short.qb and long.qb: 24h & 7d jobs on “B” nodes, 8 GB per slot
- short.qc and long.qc: 24h & 7d jobs on “C” & “D” nodes, 16 GB per slot

High-memory queues (access on request)

- himem.qh: 14d jobs on “H” nodes, 42.67 GB per slot

Test queues

- test.qc: 1 minute jobs on rescomp[1,2], 16 GB per slot

Project server queues (at the request of the owning groups)

- e.g. coolibah.q: 16 slots, no time limit, on 32-thread machine
- e.g. spencer.q: 40 slots, no time limit, on two 32-thread machines

STRUBI queues on the ResComp cluster

Electron microscopy requires running a job that mixes MPI with threading called relion

- No SGE concept of this type of job, so we need multiple special queues
- When one queue is active on a node, it suspends all the others

STRUBI relion queues

- relion1.qc: 14d MPI jobs on “C” nodes, 16 GB per slot, threading j=1
- relion2.qc: 14d MPI jobs on “C” nodes, 32 GB per slot, threading j=2
- relion4.qc: 14d MPI jobs on “C” nodes, 64 GB per slot, threading j=4
- relion8.qc: 14d MPI jobs on “C” nodes, 128 GB per slot, threading j=8
- relion16.qc: 14d MPI jobs on “C” nodes, 256 GB per slot, threading j=16

STRUBI high-memory node queue

- ginn.qh: 48 slots, no time limit, 5.33GB per slot
- Suspends running jobs on other queues on that node
- himem.qh: 42 slots, 14d jobs on “H” nodes, 42.67 GB per slot

STRUBI project server queue

- belmont.q: 32 slots, no time limit, 8 GB per slot on 32-thread machine

An SGE user

It's not enough merely to have a username on the system, you need to be defined within SGE to get queue access



“Hi, I’m Maurice!
What can I do?”

Linux username: schroff

- **Primary group:** “bioinf”
- **Additional groups:** “chimp” & “mcveang”

Storage:

- /users/mcvean/schroff (**home directory**)
- /well/bsg, /well/chimp & /well/mcvean

SGE username: schroff

- **ACL membership:** mcvean.acl
- **Default project:** mcvean.prjc
- **Additional projects:** mcvean.prja, mcvean.prjb

How to ask SGE about settings with “qconf”

- qconf –user schroff
- qconf –su mcvean.acl

Running my first cluster job: using qsub

```
[nilufer@login1 ~]$ cat script.sh
#!/bin/bash

echo "*****"
echo "Run on host: "`hostname`
echo "Operating system: "`uname -s`
echo "Username: "`whoami`
echo "Started at: "`date`
echo "*****"

R --vanilla << EOD
help()
q()
EOD

echo "*****"
echo "Finished at: "`date`
echo "*****"
exit 0
```

If we simply type: `qsub script.sh`

- Job might run, but how and where?
- Where would the output (or error) go?

All is controlled by “qsub” options

- Entered on the command line, or
- Embedded in the script in lines starting “#\$”

Basic qsub options

The shell is controlled by the first line (default is bash)

- `#!/bin/bash`

Job name should be short and useful

- `#$ -N test-job`

Project and queue (make sure that they match!)

- `#$ -P mcvean.prjc -q short.qc`

Destination for stdout and stderr

- `#$ -o stdout.log -e stderr.log -j y`
- **Defaults are** `$JOB_NAME.[o,e]$JOB_ID(.$SGE_TASK_ID)`
- **If you specify a directory, files go there with default names**

Setting the environment

- `#$ -cwd -V`
- `#$ -pe [shmem|mpi] n`

Holding a job in the queue

- `#$ -h`

Running and “array job” of separate “tasks” (`$SGE_TASK_ID`)

- `#$ -t [m-]n[:s] -tc c`

The simple script with qsub options

```
[nilufer@login1 ~]$ cat job.sh
#!/bin/bash

#$ -cwd -V
#$ -N job -j y
#$ -P zondervan.prja -q short.qa
#$ -t 1-10 -tc 2

echo "*****"
echo "SGE job ID: "$JOB_ID
echo "SGE task ID: "$SGE_TASK_ID
echo "Run on host: "`hostname`
echo "Operating system: "`uname -s`
echo "Username: "`whoami`
echo "Started at: "`date`
echo "*****"

R --vanilla << EOD
help()
q()
EOD

echo "*****"
echo "Finished at: "`date`
echo "*****"
exit 0
```

Running the simple script as a cluster job

With the “qsub” options in the script, it is simple

```
[nilufer@login1 ~]$ qsub job.sh
Your job-array 3202376.1-10:1 ("job") has been submitted
[nilufer@login1 ~]$
```

- SGE gives each job a unique job ID (and each task a unique task ID)
Job enters the scheduler queue and waits for free core(s)

mvdbunt	queued	1	208	1	26	qw=1 (110:110)
nilufer	queued	1	10	1	10	qw=1 (2700:2700)
pkalbers	queued	65	65	0	0	qw=65 (14:103)
tmills	queued	22	22	0	0	qw=22 (26:222)
vlagou	queued	1	19748	1	19748	qw=1 (2:2)

After a while a set of output files are produced

```
[nilufer@login1 ~]$ ls
job.o3202376.1  job.o3202376.2  job.o3202376.4  job.o3202376.6  job.o3202376.8
job.o3202376.10  job.o3202376.3  job.o3202376.5  job.o3202376.7  job.o3202376.9
[nilufer@login1 ~]$
```

Making the SGE job output file useful

Make the output files useful for troubleshooting

- At the “head” of file

```
[nilufer@login1 ~]$ head job.o3202376.2
*****
SGE job ID: 3202376
SGE task ID: 2
Run on host: compA000
Operating system: Linux
Username: nilufer
Started at: Thu Jun 20 09:47:26 BST 2013
*****
R version 2.14.0 (2011-10-31)
```

- At the “tail” of the file

```
[nilufer@login1 ~]$ tail job.o3202376.2

## For programmatic use:
topic <- "family"; pkg_ref <- "stats"
help((topic), (pkg_ref))

> q()
*****
Finished at: Thu Jun 20 09:47:27 BST 2013
*****
[nilufer@login1 ~]$
```

Simple reasons for a job to go wrong...

The shell script was edited using Windows

- Uses both carriage return & line feed
- User “`dos2unix -n file.in file.out`”
- Gives `execvp` error in “`qstat -j <job-ID>`”

You have the wrong shell path

- e.g. “`#!/usr/bin/csh`”
- Also gives `execvp` error in “`qstat -j <job-ID>`”

Environment not set up at submit time

- e.g. in order to use R you need to load the module for the correct version before doing “`qsub`” with the “`-V`” and “`-cwd`” options
- e.g. “`module load R/3.1.0`”

What happens when a job goes wrong?

- It can disappear from the queue, look with “`qacct -j <job-ID>`”
- It can be re-queued in error state (“`Eqw`”), look with “`qstat -j <job-ID>`”



How does SGE decide which job to run?

Your job is to submit jobs, our job is to ensure that they start

- Don't wait for free slots, don't expect jobs to start immediately!

We use the “share tree” policy

- Very flexible, but complex, includes historical usage formula
- Another area where there are known SGE bugs
- Essentially “use-it-or-lose-it” (i.e. minimal usage history)

Each group has a number of shares for each type of core

- No shares specific to any user

SGE calculates the number of running jobs for each group

For each queued job (in submission order):

- SGE assigns a priority based on how much starting that job would bring the ratios for running jobs for each group toward the ratios of shares
- if equal priorities, then jobs start first come, first served

SGE starts the jobs in priority order, repeating the scheduling calculation regularly

The share tree set-up in more detail

```
strubi.prjb.high=100
strubi.c=20
strubi.prjc.low=1
  default=1000
strubi.prjc=10
  default=1000
strubi.prjc.high=100
tomlinson.a=6
tomlinson.prja.low=1
  default=1000
tomlinson.prja=10
  default=1000
tomlinson.prja.high=100
tomlinson.b=2
tomlinson.prjb.low=1
  default=1000
tomlinson.prjb=10
  default=1000
tomlinson.prjb.high=100
tomlinson.c=30
tomlinson.prjc.low=1
  default=1000
tomlinson.prjc=10
  default=1000
tomlinson.prjc.high=100
vau.c=10
```

Display share tree using “qconf -sst”

Target share at each level of tree

- tomlinson.a=6 (shares on 288 A cores)
- tomlinson.b=2 (shares on 640 B cores)
- tomlinson.c=30 (shares on 1440 C cores)

Priority modification only within group

- tomlinson.prjc.low=1
- tomlinson.prjc=10
- tomlinson.prjc.high=100

Within a group (normal/low priority)

- Take equal turns (round robin)

Within a group (high priority)

- First come, first served

SGE control and monitoring commands

qalter: change characteristics of a queued/running job

- **Actually** qalter, qsh, qrsh, qlogin and qresub are all just qsub

qrls: release the hold on a job

qmod: clear error status on jobs and queues

qdel [-f]: delete a job

qstat: status of queued and running jobs

qacct: accounting records for completed jobs

qhost: status of execution hosts (compute nodes)

qconf: configuration of SGE

qalter: changing characteristics of a job

It is possible to change the characteristics of both queued and running jobs

Changes to running jobs may cause them to terminate and/or be re-queued!

Changes to queued jobs are usually either:

- `qalter -h u <jobid>`
- `qalter -P group.prja -q short.qa <jobid>`
- `qalter -P group.prja -q short.qa -u <myname>`

You cannot change the time limit on a running job!

- If it's not going to finish in time then kill it now (`qdel`)!

qrls: releasing the hold on a job



There are both user and system holds

You can only modify user holds

- Submit jobs with `qsub -h` or alter them with `qalter -h u <jobid>`
- Release job with `qrls <jobid>`
- Can set job start to depend on the termination of one or more existing jobs with `qsub -hold_jid <joblist>`

qmod: clearing an error state on a job

Occasionally you may be able to fix the error that makes a job go into “Eqw” state

- e.g. disk quota exceeded or if the output directory did not exist
- Clear the job error with `qmod -cj <jobid>`

This is usually a system administrators command

qdel: deleting a job



To delete running or queued jobs use qdel

- qdel <jobid>
- qdel -f <jobid> if node is dead or for stuck MPI jobs

qstat: checking the status of jobs



Report on the status of just my jobs with qstat

- Use `qstat -s p` for pending jobs and `qstat -s r` for running jobs

Report on the status of all jobs in the cluster with `qstat -u "*"`

- Output can be very long and indigestible

```
[nilufer@login1 ~]$ qstat -u ashish
```

job-ID	prior	name	user	state	submit/start at	queue	slots	ja-task-ID
3202368	541.66672	R4_baserec	ashish	r	06/20/2013 05:55:45	gpfs.qb@compB033.cluster3	2 1	
3202368	541.66672	R4_baserec	ashish	r	06/20/2013 05:55:45	gpfs.qb@compB033.cluster3	2 3	
3202368	541.66672	R4_baserec	ashish	r	06/20/2013 05:55:45	gpfs.qb@compB033.cluster3	2 4	
3202369	0.00000	R4_printre	ashish	hqw	06/20/2013 01:19:55		2 1-20:1	
3202370	0.00000	R4_reducer	ashish	hqw	06/20/2013 01:19:55		2 1-20:1	

Report detailed information on a job with `qstat -j <job_id>`

- Output also long and indigestible

```
script_file: /gpfs1/well/projects8/gotzd/SUMMIT/Iund/BamPr/bin4/a.10.base_recalibratorKUN4.sh
parallel environment: shmem range: 2
jid_predecessor_list (req): R4_indexdedup20bams
jid_successor_list: 3202369
verify_suitable_queues: 2
project: gpfs.prjb
job-array tasks: 1-20:1
usage 1: cpu=00:08:36, mem=189.41702 GBs, io=0.39162, vmem=441.582M, maxvmem=441.582M
usage 3: cpu=02:33:32, mem=3397.68245 GBs, io=6.01999, vmem=443.020M, maxvmem=443.020M
usage 4: cpu=04:35:35, mem=7108.02473 GBs, io=12.96924, vmem=505.566M, maxvmem=505.566M
scheduling info: queue instance "long.qb@compB017.cluster3" dropped because it is temporarily not available
queue instance "short.qb@compB017.cluster3" dropped because it is temporarily not available
queue instance "short.qb@compB004.cluster3" is in suspend alarm: load short=14 280000 (no 1s
```

qacct: accounting of finished jobs



When jobs have finished qstat forgets about them

```
[nilufer@login1 ~]$ qstat -j 3202376
Following jobs do not exist:
3202376
[nilufer@login1 ~]$
```

Some records remain accessible through qacct

Aggregated accounting information e.g. for back charging

- Andrew Morris has used just over 25 CPU/years
- We do not back charge for actual use.
- Can take a long time to run!
- Can produce a large amount of output!

```
[nilufer@login1 ~]$ qacct -o amorris
```

OWNER	WALLCLOCK	UTIME	STIME	CPU	MEMORY	IO
amorris	837830920	713522679.978	707494.717	781348256.501	17751591862.125	252805.412

```
[nilufer@login1 ~]$
```

You can also get more detailed information on a specific job(s)...

qacct: accounting of a single job or task

Use `qacct -j <job_id> [-t <task>]`

- qname: **queue** job ran in
- hostname: **where** job ran
- jobnumber: **job ID**
- taskid: **number of task** for array job
- start_time: **when** job started
- end_time: **when** job ended
- failed: **should be 0**
- exit_status: **exit status of script:**
 - <128 user defined
 - >128 SGE signal
- 134 = SIGABRT e.g. memory exceeded
- 137 = SIGKILL e.g. time exceeded
- ru_wallclock: **elapsed time**
- ru_utime: **user cpu time**
- cpu: **cpu time used**
- maxvmem: **maximum size reached by job**

```
[nilufer@login1 ~]$ qacct -j 3202376 -t 2
```

```
=====
qname          short.qa
hostname       compA000.cluster3
group          zondervan
owner          nilufer
project        zondervan.prja
department     defaultdepartment
jobname        job
jobnumber      3202376
taskid         2
account        sge
priority       0
qsub_time      Thu Jun 20 09:46:34 2013
start_time     Thu Jun 20 09:47:26 2013
end_time       Thu Jun 20 09:47:27 2013
granted_pe     NONE
slots          1
failed         0
exit_status    0
ru_wallclock   1
ru_utime       0.322
ru_stime       0.058
ru_maxrss     0
ru_ixrss      0
ru_ismrss     0
ru_idrss      0
ru_isrss      0
ru_minflt     13165
ru_majflt     0
ru_nswap      0
ru_inblock    0
ru_oublock    0
ru_msgsnd     0
ru_msgrcv     0
ru_nsignals   0
ru_nvcsw     549
ru_nivcsw     32
cpu           0.380
mem           0.000
io            0.000
iow           0.000
maxvmem       0.000
```



qghost and qconf: housekeeping tools



qghost: reports on the state of execution hosts

- Looking for dead nodes,
- Reporting how many cores are in use
- Reporting how much memory (and swap) are in use
- Checking state of per-node queue instances e.g. qghost -q

qconf: sets and reports on all the administrative setup of SGE

- `qconf -[s|a|m|M|d]<feature>[l| <name>| <filename>]`

`qconf -s<feature>l`: list all feature elements

- `qconf -suserl`
- `qconf -sql`

`qconf -s<feature> <name>` reports one element of a feature

- `qconf -suser nilufer`
- `qconf -sq short.qa`
- `qconf -shgrp @short.hga`

... and now for something a bit simpler

The output of qstat can be seriously verbose!

- We have two in-house monitoring tools:

qsum (see qsum -h)

- reformatted from the output of qstat -ext -u "*"
- compactly summarizes running and queueing jobs

qload (see qload -h, qload -ho & qload -hn)

- reformatted from the outputs of qstat -s rs -ext -t -u "*" and qhost -q
- simple text-based diagram of running jobs

Both probably still have bugs in!

Thank you for your attention... any questions?
Next talk 11:00 on Monday... "Monitoring and troubleshooting"

