

The WTCHG Research Computing Core

Talk 4: Monitoring and troubleshooting

Robert Esnouf (robert@strubi.ox.ac.uk; robert@well.ox.ac.uk),
Jon Diprose (jon@well.ox.ac.uk) & Colin Freeman (cfreeman@well.ox.ac.uk)

Generic emails: support rescomp@well.ox.ac.uk
users rescomp-users@well.ox.ac.uk

A series of introductory talks...

A set of six talks of about 45 minutes each (plus questions)

• Talk 1: What is the ResComp Core?

(Mon 23/1 10:00 Room B; Robert Esnouf)

Talk 2: A basic introduction to Linux

(Wed 25/1 10:00 Room B; Robert Esnouf)

Talk 3: Submitting jobs to the cluster

(Thu 26/1 10:00 Room B; Robert Esnouf)

Talk 4: Monitoring and troubleshooting

(Mon 30/1 11:00 Room B; Robert Esnouf)

Talk 5: ResComp centrally-managed applications

(Wed 1/2 10:00 Room B; Jon Diprose)

Talk 6: Doing your own thing (compiling and customizing)

(Thu 2/2 10:00 Room B; Jon Diprose)

The story so far...

The ResComp cluster has different types of node:

- 288 “A” cores @ 2.67GB/core; 640 “B” cores @ 8GB/core; 1720 “C” cores & 768 “D” cores @ 16GB/core; 144 “H” cores @ 42.66GB/core
- Login/submission nodes: rescomp[1-2] and the project servers

Files must be on the (chargeable) GPFS filesystems:

- /gpfs0: */users/<group>/<user>*, */apps/well*, */mgmt*
- /gpfs1 & /gpfs2: */well/<group>*

Refresher of basic Linux concepts

- User accounts, groups, commands, shells (bash)
- Environments and environment variables: \$PATH and \$LD_LIBRARY_PATH
- File streams (stdin, stdout and stderr) redirection & pipes (“<”, “>”, “|”)
- Foreground and background jobs (“&”), “nohup” and “screen”
- Writing, executing and sourcing a shell script

The story so far...

Cluster computing is non-graphical and non-interactive!

- Sun Grid Engine (SGE) using a group-based share tree policy

Your job is to submit jobs, our job is to ensure that they start

- Don't wait for free slots, don't expect jobs to start immediately!

Submit scripts to the SGE scheduler using “qsub”

- Embed “qsub” arguments using “#\$” lines
- Specify where stdout and stderr should go
- “#\$ -pe shmem|mpi <n>” for multicore or MPI jobs
- Modify jobs with “qalter”, “qrls”, “qmod” & “qdel”
- Monitor jobs and SGE with “qstat”, “qacct”, “qhost” & “qconf”

Different queues: typically short (<24 hours) & long (<7 days)

- Jobs killed if they exceed time or memory limits
- Special queues for relion jobs and on project servers
- Use “qstat -j <job_id>” for details on queued and running jobs
- Use “qacct -j <job_id> [-t task]” for details of completed jobs

Overview of this talk...

Looking at “qstat” in more detail

The in-house monitoring commands

“qsum”

“qload”

A reminder of “qacct” for finished jobs

The ResComp Cluster workshop

Looking at “qstat” in more detail

```
-bash-4.1$ qstat -u hatton
job-ID prior   name       user          state submit/start at   queue                          slots ja-task-ID
-----
1979389 0.20626 POBI_GATK_ hatton      r    01/21/2015 22:24:16 himem.qh@compH000             16
1979390 0.00000 POBI_GATK_ hatton      hqw   01/21/2015 22:27:00                               30
1979391 0.00000 POBI_GATK_ hatton      hqw   01/21/2015 22:28:50                               30
-bash-4.1$
```

“qstat” is main tool for looking at queued and running jobs

- By default “qstat” just shows your jobs
- use “qstat -u <user,user...>” for other people
- use “qstat -u “*”” for everyone
- Add “-s rs” for running/suspended jobs, “-s p” for pending jobs

Queued jobs sorted by priority, fields are:

- \$JOB_ID, priority, \$NAME, \$USER, state, submit or start time, execution queue instance, # slots, [\$SGE_TASK_ID or task range]

Output is typically 2,000–20,000 lines long for all users!

“qstat” job state strings

```
-bash-4.1$ qstat -u hatton
job-ID prior name user state submit/start at queue slots ja-task-ID
-----
1979389 0.20626 POBI_GATK_ hatton r 01/21/2015 22:24:16 himem.qh@compH000 16
1979390 0.00000 POBI_GATK_ hatton hqw 01/21/2015 22:27:00 30
1979391 0.00000 POBI_GATK_ hatton hqw 01/21/2015 22:28:50 30
-bash-4.1$
```

Running jobs

- r: running
- t: transitioning (in or out, usually in)
- dr: registered for deletion (e.g. a dead node; “qdel -f <job_id>”)

Suspended jobs (no user action required)

- s: suspended by a superior queue (queue subordination)
- S: suspended by administrator
- T: threshold suspension

Queued jobs

- qw: queue wait
- hqw: held in queue wait
- Eqw: error state held in queue wait (“qmod -cj <job_id>”)

Job suspension...

A suspended job is one that just sits there:

- It has an inactive process (and shepherd)
- It consumes resources such as memory
- It's wall time still ticks although no work is being done

Jobs suspended buy a superior queue (state “s”)

- Only by `ginn.qh@compH002` at the moment
- Not considered generally useful as memory is committed anyway
- Pushing a process memory into swap space is s..l..o..w..

Jobs suspended buy a an administrator (state “S”)

- We may suspend either a job or a queue without warning
- Fairly rare event, usually across the whole cluster
- Main use is to reduce I/O load temporarily, or perhaps power load

Jobs suspended owing to threshold limits (state “T”)

- The `execd` monitors (“uptime”) load on each execution host
- If short-time load exceeds $(\#cores+2)$ no new jobs start
- If short-time load exceeds $(\#cores \times 1.5)$ jobs suspended “at random”

Threshold suspension...

If short-time load exceeds ($\text{\#cores} \times 1.5$) jobs are suspended

- The suspended job (state “T”) may not be the one causing the stress
- Ask for >1 slot with “ $\$ -pe \text{ shmem } <n>$ ” and $\$NSLOTS$ is set for you
- $\$NSLOTS$ is not set (to be 1) by default

Clusters are very different to workstations

- Workstations you often only use one core, rest are mostly idle
- Any parallelization for a workstation is usually good
- Clusters use schedulers to pack jobs onto all cores at all times
- Imperfect parallelization is often bad(!)
- For a job set the first job might finish earlier, last job will finish later

SGE will not play nicely with some modern parallelizations

- “pthreads” cause particular problems
- Fast context switching: great for high-latency asynchronous processes
- Classic example is network related activity, often just waiting for work
- Need to run of special queues without threshold suspension
- For mathematical computing benefits seem modest

... and now for something a bit simpler

The output of `qstat` can be seriously verbose

- Output is typically 2,000–20,000 lines long for all users!

We have two in-house monitoring tools:

`qsum` (see `qsum -h`)

- reformatted from the output of `qstat -ext -u "*"`
- compactly summarizes running and queueing jobs

`qload` (see `qload -h`, `qload -ho` & `qload -hn`)

- reformatted from the outputs of `qstat -s rs -ext -t -u "*"` and `qhost -q`
- simple text-based diagram of running jobs
- `qcheck` is an alias for "`qload -nall -wn 1 -xp`"

Both probably still have bugs in!



A summary of running & queuing jobs

```
-bash-4.1$ qsum
```

USERS	QUEUES	JOBS	CORES	ARRYS	TASKS	PROJECTS	STATES (FIRST:LAST)
abhay	relion.qc	2	192	0	0	strubi.prjc	r=2
acortes	short.qc	160	160	0	0	mcvean.prjc	r=160
amorris	long.qc	593	593	0	0	morris.prjc	r=593
amorris	short.qc	8	64	0	0	morris.prjc	r=8
batty	short.qc	297	306	21	296	donnelly.prjc	r=297
frot	long.qc	67	67	0	0	mcvean.prjc	r=67
hatton	himem.qh	1	16	0	0	donnelly.prjc	r=1
juha	belmont.q	1	14	0	0	strubi.prjc	r=1
kvik	short.qc	1	1	0	0	bsg.prjc	r=1
leffler	spencer.q	1	1	0	0	spencer.prjc	r=1
luigi	relion.qc	1	96	0	0	strubi.prjc	r=1
petkova	long.qc	1	1	1	1	donnelly.prjc	dr=1
schroff	long.qc	21	126	0	0	mcvean.prjc	r=21
sknight	short.qc	1	1	0	0	brc.prjc	r=1
tomparks	long.qc	2	2	0	0	hill.prjc	r=2
RUNNING		1157	1640	22	297		
abhay	queued	1	24	0	0	strubi.prjc	qw=1 (4:4)
acortes	queued	872	872	0	0	mcvean.prjc	qw=872 (595:2749)
amorris	queued	241	1928	0	0	morris.prjc	qw=241 (792:1466)
batty	queued	403	8059	403	8059	donnelly.prjc	qw=403 (1578:3135)
frot	queued	907	907	0	0	mcvean.prjc	qw=907 (39:1725)
hatton	held	2	60	0	0	donnelly.prjc	hqw=2 (3136:3137)
ip	queued	1	10	0	0	bsg.prjc	qw=1 (1:1)
juha	queued	1	24	0	0	strubi.prjc	qw=1 (2:2)
luigi	queued	1	24	0	0	strubi.prjc	qw=1 (3:3)
schroff	queued	708	4248	0	0	mcvean.prjc	qw=708 (5:2370)
QUEUED		3137	16156	403	8059		

qsum: what about me?

You can ask qsum just to report on named user(s)

- `qsum -u <username,username...>`
- `qsum -u "*"`

If you omit the username, qsum assumes you are interested in yourself

- `qsum -u`
- Equivalent to a compact summary of the output of `qstat`

```
-bash-4.1$ qsum -u
```

USERS	QUEUES	JOBS	CORES	ARRYS	TASKS	PROJECTS	STATES (FIRST:LAST)
frot	long.qc	71	71	0	0	mcvean.prjc	r=71
	RUNNING	71	71	0	0		
frot	queued	907	907	0	0	mcvean.prjc	qw=907 (1:907)
	QUEUED	907	907	0	0		

qsum: what about us?

You can just ask qsum to report on a named project(s)

- `qsum -p <proj,proj...>`
- Matches any project name for which each *proj* is a substring
- e.g. `mcvean` matches `mcvean.prja`, `mcvean.prjb` & `mcvean.prjc`
- e.g. `prjb` matches all groups `group.prjb` projects

If you omit the project name, qsum assumes you are interested in your default project base name

- `qsum -p`

```
-bash-4.1$ qsum -p
```

USERS	QUEUES	JOBS	CORES	ARRYS	TASKS	PROJECTS	STATES (FIRST:LAST)
acortes	short.qc	158	158	0	0	mcvean.prjc	r=158
frot	long.qc	71	71	0	0	mcvean.prjc	r=71
schroff	long.qc	21	126	0	0	mcvean.prjc	r=21
RUNNING		250	355	0	0		
acortes	queued	875	875	0	0	mcvean.prjc	qw=875 (692:2490)
frot	queued	907	907	0	0	mcvean.prjc	qw=907 (29:1374)
schroff	queued	708	4248	0	0	mcvean.prjc	qw=708 (1:2082)
QUEUED		2490	6030	0	0		

Why is no queue listed for queued jobs?

qsum information is farmed from the output of qstat

- qstat does not show queue information for queued jobs

A job can be submitted to many queues

- `qsub -q short.qa,short.qb,short.qc -P mcvean.prja <script.sh>`

Knowing about target queues could be very useful

- Information available in `qstat -j <job_id>`
- `qsum -Q` (limited to first 100 jobs in the queue)

```
-bash-4.1$ qsum -Q -u amorris
*** Too many queued jobs for "-q"
*** Remaining jobs listed simply as "queued"
```

USERS	QUEUES	JOBS	CORES	ARRYS	TASKS	PROJECTS	STATES (FIRST:LAST)
amorris	long.qc	590	590	0	0	morris.prjc	r=590
amorris	short.qc	8	64	0	0	morris.prjc	r=8
RUNNING		598	654	0	0		
amorris	queued	141	1128	0	0	morris.prjc	qw=141 (101:241)
amorris	short.qc	100	800	0	0	morris.prjc	qw=100 (1:100)
QUEUED		241	1928	0	0		

What else can qsum do?

Not much else at the moment!

- Check with `qsum -h`

```
-bash-4.1$ qsum -h
```

qsum is a script to show the ownership of running and queued jobs across the cluster. The number of cores for each queue slot is defined in `qqueue.conf` in the same directory as `qsum`. With no arguments the output is for all users.

Syntax: `qsum [-h] [-u [USER,...]] [-p [PROJECT,...]] [-q]`

<code>-h</code>	- show this help information
<code>-u [USER,...]</code>	- shows jobs for the specified user(s) only If no USER then then current user is shown
<code>-p [PROJECT,...]</code>	- shows jobs for the specified project(s) only If no PROJECT then the current user's default project is used (matches all projects with PROJECT as any substring)
<code>-Q</code>	- shows target queues for submitted jobs [limited to 100 queued jobs unless "root"]

One newer feature is `qsum -P`

- Just displays “pending” (i.e. queueing and held) jobs

What are the compute nodes doing?

The in-house command to look at this is qload

- qload has many arguments, but for most users you won't need any
- qload -h: describes the input arguments
- qload -hn: describes how to select groups of nodes
- qload -ho: describes the output display

```
[root@head1 ~]# qload
JOBS FOR SELECTED NODES:
compA000 *****
compA009 *****
compA018 *****
compA027 *****
compB000 *****
compB009 *****
compB018 *****
compB027 *****
compB036 *****

compC000 L>LLLLLLLLLLLLLLLL L>LLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL L>LLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL
compC007 LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL L>LLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL
compC014 LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL
compC021 LLLLLLLLLLLLLLLLL L>L>LLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL
compC028 LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL
compC035 L>LLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL ***** LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL
compC042 LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL LLLLLLLLLLLLLLLLL L>LLLLLLLLLLLLLLLL SSSSSSSSSSSSSSSSS
compC049 SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS
compC056 SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS
compC063 SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS
compC070 SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS
compC077 SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS
compC084 SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS SSSSSSSSSSSSSSSSS
compC091 ..... f-->f-->f-->f--> f-->f-->f-->f--> ..... f-->f-->f-->f--> .....
compC098 ..... f-->f-->f-->f--> f-->f-->f-->f--> ..... f-->f-->f-->f--> .....
compC105 ..... R-->r-->r-->r--> f-->f-->f-->f--> ..... f-->f-->f-->f--> .....

[root@head1 ~]#
```

Basic output of qload

The default view is to show the main cluster nodes

- The view is automatically adjusted to the terminal width

Multiple nodes are shown on one line (normally)

Only the first node on each line is named

- You can usually just count along

Nodes are separated by horizontal bars depending on number of slots (cores)

Each node is represented by a string of characters that is one more than the number of cores it offers

- Most characters are what each core is doing (next slide)
- The last character is a status code, usually blank for a happy node

What is each core doing?

There is one character per core for each node

- “.”: A free core on which jobs can be scheduled
- “:”: A core in use, but not for a job specified by input flags
- “,”: A core on which jobs cannot be scheduled in the current config
- “*”: A core on a node which is not running an execd (usually dead)
- “L”: A job running in a long queue
- “S”: A job running in a short queue
- “R”: A job running in a relion queue
- “X”: A job running in a special queue
- “L”, “L>”, “L->”, “L-->”: A job occupying 1, 2, 3 or 4 cores
- “r-->”: A lower case letter is a slave process in multi-node MPI jobs

```
[root@head1 ~]# qcheck
JOBS FOR SELECTED NODES:
-----+-----
compC010 | LLLLSSSSSSSSSSLLLLd
compC037 | *****
compC067 | L---->SSSSSSSSSSLo
-----+-----
```


What are the status codes?

There can be many possible issues with a node

- A single letter represents the most serious issue with the node (last listed issue)
- “ ”: A blank status code indicates a happy node (the usual state)
- One of “aoACD”: usually a temporary issue (load threshold, orphaned, suspend threshold, calendar suspend, calendar disable)
- “<”: A node under stress: *i.e.* high load, memory use or swap use
- One of “cdsuE”: usually an issue requiring an administrator (config ambiguous, disabled queue, suspended queue, unknown state, error)
- “+”: A node on which there are one or more suspended jobs
- “*”: A node which is not running an execd (usually dead)

```
[root@head1 ~]# qcheck
JOBS FOR SELECTED NODES:
-----+-----
compC010 | LLLLSSSSSSSSSSLLd
compC037 | *****
compC067 | L---->SSSSSSSSLo
-----+-----
```

How stressed are the nodes?

You can ask qload to summarize the stress on each node

- Use the “-l” option to display load and memory usage (in GB)

```
jaga      | ..... 2.1 8G
liono     | X..... 0.1 4G
jarrah    | ..... 0.3 4G
belmont   | X----->..... 11.7 15G
```

- Use the “-ls” option to additionally display swap usage (in GB)

```
jaga      | ..... 2.1 8G 0S
liono     | X..... 0.1 4G 0S
jarrah    | ..... 0.3 4G 0S
belmont   | X----->..... 11.7 15G 0S
```

- Swap usage is usually zero, so not interesting!

Selecting sets of nodes with qload

qload has a very flexible definition of which nodes to display

- qload defaults to showing the main (A, B, C & D) compute nodes
- qload -np [-wn 1] shows just the project servers
- qload -nh shows just the high-memory servers
- qload -n belmont or qload belmont shows just one node (in detail)

This is usually simple to use and probably best demonstrated by some examples:

jaga,liono	Nodes named jaga and liono
compC[000-107]	Nodes named from compC000 to compC107
%a	Nodes named in the group %a, in this case the "A" nodes: nodes from compA000 to compA035
%b,^compB017	Nodes in the group %b, in this case the "B" nodes, but node compB017 is excluded from the list
%bc6	Nodes in the group %bc6, in this case blade chassis number 6: nodes from compC090 to compC107
compC[000-107:18]	Nodes in slot 1 in each blade chassis
@relion.hgc	Nodes in the SGE hostgroup @relion.hgc, in this case nodes that can run "relion" MPI jobs

Currently defined GROUPS are: %all - all managed nodes; %ab - A and B nodes; %abc - A, B and C nodes; %a - A nodes; %b - B nodes; %c - C nodes; %h - high memory nodes; %p - the project servers; %t - the rescomp nodes; %bc[n] - the nodes in each Fujitsu blade chassis, from 1 to 6

The current default selection is equivalent to "-n %abc"

Looking at the load on a single node

If you select a single node qload gives more detailed output

- Detailed output is also turned on with the “verbose” flag “-v”
- `qload -nh -v` shows just the detailed load on the high-memory servers

```
[root@head1 ~]# qload -nh -v
DETAILED QUEUE, JOB, MEMORY AND SWAP USE FOR ALL SELECTED NODES:
-----+-----
compH000 | H-----> 48.0   22G   0S
-----+-----
  himem.qh          BIP   0/48/48
-----+-----
8857370   0.000 segmentref juha          strubi.prjc    r    01/20/2016 10:50:25 himem.qh@compH000          48
-----+-----
compH001 | R-----> 16.9   219G   0Sd
-----+-----
  himem.qh          BIP   0/0/48          d
  relion.qh         BIP   0/16/16
-----+-----
8851461   0.092 relion      mgrange      strubi.prjc    r    01/20/2016 11:40:16 relion.qh@compH001          48
-----+-----
compH002 | G-----> 48.0   21G   0S+
-----+-----
  himem.qh          BIP   0/0/42          dS
  ginn.qh           BIP   0/48/48
  long.qc           BIP   0/6/48          So
-----+-----
8855637   0.008 fm          asanniti      jknight.prjc    S    01/19/2016 13:10:51 suspended@compH002          0
8855793   0.092 index       ginn          strubi.prjc    r    01/19/2016 15:37:43 ginn.qh@compH002          48
-----+-----
```

Other command line options for qload

qload has lots of useful command line options, but mostly

aimed at system administrators

```
Syntax:  qload [-h|-hn|-ho] [-j JOB_ID,...] [-u [USER,...]] [-p [PROJ,...]]
          [-q [QUEUE,...]] [-l[s]] [-d] [-v] [-x] [-xp]
          [-wc CHARS|-wn NODES] [-nGROUP|[-n] LIST]
```

-h	- show this help information
-hn	- show help information for node list and group selection
-ho	- show help information for the output format
-j JOB_ID,...	- explicitly show the specified job ID(s) (overrides -u, -p and -q)
-u [USER,...]	- explicitly show jobs for the specified user(s) only (not with -j)
-p [PROJ,...]	- explicitly show jobs for the specified project(s) only (not with -j) (matches all projects with PROJ as any substring)
-q [QUEUE,...]	- explicitly show jobs for the specified queue(s) only (not with -j)
-l[s]	- show load, memory use and [with ls] swap use in addition to jobs
-d	- show information only for selected "dead" nodes)
-v	- show verbose information for all selected nodes (default for a single node)
-x	- show exclusively the subset of nodes running (selected) jobs
-xp	- show exclusively the subset of nodes currently reporting problems
-wc CHARS	- width of line in characters (only one of -wc or -wn)
-wn NODES	- width of line in nodes (only one of -wc or -wn)
-nGROUP	- show information for the specified group of nodes
[-n] LIST	- show information for the specified list of nodes

What are the project servers doing?

The project servers are not shown in default qload output

- Select the node group “p” for the known project servers
- Select one node per line for the output
- Ask for details on the load, memory and swap space usage
- `qload -np -wn 1 -ls`

```
[root@head1 ~]# qload -np -wn 1 -ls
JOBS, LOAD, MEMORY AND SWAP USE FOR SELECTED NODES:
-----+-----
jaga      | ..... 2.1  8G  0S
liono     | X..... 0.1  4G  0S
jarrah    | ..... 0.3  4G  0S
belmont   | X----->..... 11.7 15G 0S
-----+-----
coolibah  | ..... 0.0  4G  0S
-----+-----
```

Are there any obvious problems?

qcheck reports on obvious potential cluster issues

- Mainly intended for system administrators
- qcheck is an alias for “qload -nall -wn 1 -xp ”

```
[root@head1 ~]# qcheck
JOBS FOR SELECTED NODES:
-----+-----
compC010 | LLLLSSSSSSSSSSLLd
compC037 | *****
compC067 | L---->SSSSSSSSLo
-----+-----
```

“qacct” holds all records of cluster activity

Last year, Winni was winning the race: 2,872,418,832s = 91y

- Note: WALLCLOCK can be greater than or less than CPU

```
[root@head1 ~]# head -n 2 qacct.log; qacct -o | sort -n -r -k 5
```

OWNER	WALLCLOCK	UTIME	STIME	CPU	MEMORY	IO
winni	2382301270	2484037020.560	86901869.432	2872418832.182	62650812919.367	341297.179
abhay	238951659	1777352696.697	5636305.695	2442072111.750	7342403576.611	555685.395
quang	2542659858	1409670121.162	15253423.546	1957784086.727	6076923543.409	176989781.903
pkalbers	1142779152	805972918.808	7981886.165	1221543516.007	4178093461.965	28633874.572
frot	947906037	662813387.896	550379.697	942510221.023	997703349.339	153688.228
joezhu	824565186	340958934.266	381637.493	790850679.743	544538501.170	53450.581
amorris	690359156	623190114.065	761945.105	682648343.852	6579112255.301	62643.949
jfertaj	596822976	152187281.841	140985.078	577939910.148	2622039611.389	1418.408
oliver	478883764	466395403.787	3481083.429	472707510.995	86169920.099	23928193.913
schroff	438766689	371283525.901	413322.945	437491249.252	4795068604.388	38245.680
caina	451845868	399797873.482	9525394.522	429257863.259	2862026844.662	806922.898
batty	393329336	364774976.372	5316687.214	409486349.341	2763569699.786	938866.945
rwDavies	336233233	275059782.512	232214.047	332711743.168	654071255.454	286641.673
marchini	323791393	258724828.564	234473.072	320104407.769	12955558.847	4635.516
kelliott	590088552	285095615.337	7203728.199	304268302.842	2390072557.185	12387104.549
hatton	108858489	291071575.582	2195420.615	302895720.820	3185567938.943	1665207.424
vivekn	76437435	221658629.886	3236468.829	286479285.941	498629314.345	40618.322
george	361469236	268367822.955	1215984.700	272877916.849	77843980.304	3594082.408
clare	252384384	243803011.698	538274.621	245080599.465	236638233.951	792989.495
hiliary	167915762	205278037.913	2691833.167	229413600.656	488251023.789	103563.557
juha	35572267	149301187.925	1131980.848	220528803.910	762496625.647	374600.590
gav	190481424	179369941.281	4385933.474	185636924.738	397727783.982	5075.299
acortes	179624039	172528141.517	398948.398	177127313.095	349164778.034	154049.526
davismcc	72767118	164915378.712	1376505.926	167866374.083	115951363.674	47067.287
rivas	196181952	141741895.327	1935606.141	150215822.117	34648132505.392	301986.678
mvdbunt	114634728	114921614.027	20386271.512	137896111.279	888692965.572	866882.327
leffler	123206138	96081001.791	33710361.695	134097940.508	78507748.458	2173864.879
wadh4116	54687282	103261708.562	400207.291	131159344.302	1804498138.848	288616.355

“qacct” for a single job or task

Use `qacct -j <job_id> [-t <task>]`

- qname: **queue** job ran in
- hostname: **where** job ran
- jobnumber: **job ID**
- taskid: **number of task** for array job
- start_time: **when** job started
- end_time: **when** job ended
- failed: **should be 0**
- exit_status: **exit status of script:**
 - <128 user defined
 - >128 SGE signal
- 134 = SIGABRT e.g. memory exceeded
- 137 = SIGKILL e.g. time exceeded
- ru_wallclock: **elapsed time**
- ru_utime: **user cpu time**
- cpu: **cpu time used**
- maxvmem: **maximum size reached by job**

```
[nilufer@login1 ~]$ qacct -j 3202376 -t 2
```

```
=====
qname          short.qa
hostname       compA000.cluster3
group          zondervan
owner          nilufer
project        zondervan.prja
department     defaultdepartment
jobname        job
jobnumber      3202376
taskid         2
account        sge
priority       0
qsub_time      Thu Jun 20 09:46:34 2013
start_time     Thu Jun 20 09:47:26 2013
end_time       Thu Jun 20 09:47:27 2013
granted_pe     NONE
slots          1
failed         0
exit_status    0
ru_wallclock   1
ru_utime       0.322
ru_stime       0.058
ru_maxrss     0
ru_ixrss      0
ru_ismrss     0
ru_idrss      0
ru_isrss      0
ru_minflt     13165
ru_majflt     0
ru_nswap      0
ru_inblock    0
ru_oublock    0
ru_msgsnd     0
ru_msgrcv     0
ru_nsignals   0
ru_nvcsw     549
ru_nivcsw     32
cpu           0.380
mem           0.000
io            0.000
iow           0.000
maxvmem       0.000
```



... that's enough from me!

The final two talks are given by Jon Diprose:

- Talk 5: ResComp centrally-managed applications
(1/2 10:00 Room B; Jon Diprose)

Talk 6: Doing your own thing (compiling and customizing)
(2/2 10:00 Room B; Jon Diprose)

Thank you for your attention... any questions?
Next talk 10:00am Wednesday... "ResComp centrally-managed applications"