

The WTCHG Research Computing Core

Talk 5: The ResComp Managed Applications

Robert Esnouf (robert@strubi.ox.ac.uk; robert@well.ox.ac.uk),
Jon Diprose (jon@well.ox.ac.uk) & Colin Freeman (cfreeman@well.ox.ac.uk)

Generic emails: support rescomp@well.ox.ac.uk
users rescomp-users@well.ox.ac.uk

A series of introductory talks...

A set of six talks of about 45 minutes each (plus questions)

• Talk 1: What is the ResComp Core?

(Mon 23/1 10:00 Room B; Robert Esnouf)

Talk 2: A basic introduction to Linux

(Wed 25/1 10:00 Room B; Robert Esnouf)

Talk 3: Submitting jobs to the cluster

(Thu 26/1 10:00 Room B; Robert Esnouf)

Talk 4: Monitoring and troubleshooting

(Mon 30/1 11:00 Room B; Robert Esnouf)

Talk 5: ResComp centrally-managed applications

(Wed 1/2 10:00 Room B; Jon Diprose)

Talk 6: Doing your own thing (compiling and customizing)

(Thu 2/2 10:00 Room B; Jon Diprose)

The story so far...

The ResComp cluster has different types of node:

- 288 “A” cores @ 2.67GB/core; 640 “B” cores @ 8GB/core; 1720 “C” cores & 768 “D” cores @ 16GB/core; 144 “H” cores @ 42.66GB/core
- Login/submission nodes: rescomp[1-2] and the project servers

Files must be on the (chargeable) GPFS filesystems:

- /gpfs0: */users/<group>/<user>*, */apps/well*, */mgmt*
- /gpfs1 & /gpfs2: */well/<group>*

Refresher of basic Linux concepts

- User accounts, groups, commands, shells (bash)
- Environments and environment variables: \$PATH and \$LD_LIBRARY_PATH
- File streams (stdin, stdout and stderr) redirection & pipes (“<”, “>”, “|”)
- Foreground and background jobs (“&”), “nohup” and “screen”
- Writing, executing and sourcing a shell script

The story so far...

Cluster computing is non-graphical and non-interactive!

- Sun Grid Engine (SGE) using a group-based share tree policy

Your job is to submit jobs, our job is to ensure that they start

- Don't wait for free slots, don't expect jobs to start immediately!

Submit scripts to the SGE scheduler using “qsub”

- Embed “qsub” arguments using “#\$” lines
- Specify where stdout and stderr should go
- “#\$ -pe shmem|mpi <n>” for multicore or MPI jobs
- Modify jobs with “qalter”, “qrls”, “qmod” & “qdel”
- Monitor jobs and SGE with “qstat”, “qsum”, “qacct”, “qload” & “qconf”

Different queues: typically short (<24 hours) & long (<7 days)

- Jobs killed if they exceed time or memory limits
- Special queues for relion jobs and on project servers
- Use “qstat -j <job_id>” for details on queued and running jobs
- Use “qacct -j <job_id> [-t task]” for details of completed jobs

Overview of this talk...

- Where to find the applications

Different types of “application”

Relevant environment variables

Setting up your environment

The challenges of modular applications

Diagnosing failures

Where To Find The Applications

/apps/well/

- possibly all some of you need to know...
- 253 different applications/libraries

/apps/strubi

- don't ask me why relion is under /apps/well

/apps/htseq

- applications used by the Sequencing Core pipelines

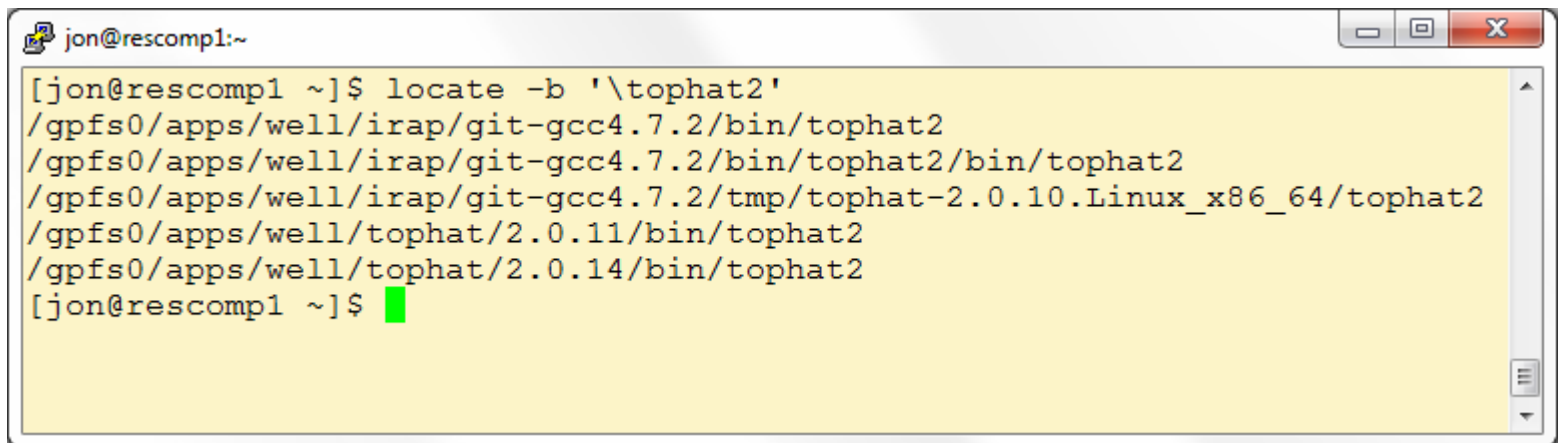
Why can't I just type, e.g., `tophat2`?

- the application version matters
- you can, but you'll need to do a bit of work...

Where To Find The Applications

- The 'locate' command on rescomp1/rescomp2 can be used to find a file

For an exact match use, eg, 'locate -b 'tophat2''



```
jon@rescomp1:~$ locate -b 'tophat2'
/gpfs0/apps/well/irap/git-gcc4.7.2/bin/tophat2
/gpfs0/apps/well/irap/git-gcc4.7.2/bin/tophat2/bin/tophat2
/gpfs0/apps/well/irap/git-gcc4.7.2/tmp/tophat-2.0.10.Linux_x86_64/tophat2
/gpfs0/apps/well/tophat/2.0.11/bin/tophat2
/gpfs0/apps/well/tophat/2.0.14/bin/tophat2
jon@rescomp1 ~]$
```


Different types of “application”

Binary

- produced by compilation of source code
- hard to understand and modify without source
- static
 - “no” dependencies – but may be making assumptions
- dynamic
 - depends on binary libraries that are not part of this binary

Script

- source code that is “interpreted” at runtime
- interpreter specified by first line starting #!
- easier to understand and to modify
- python, R, perl, shell (sh/bash/csh/tcsh/zsh) & env

The command ``file`` can tell the difference

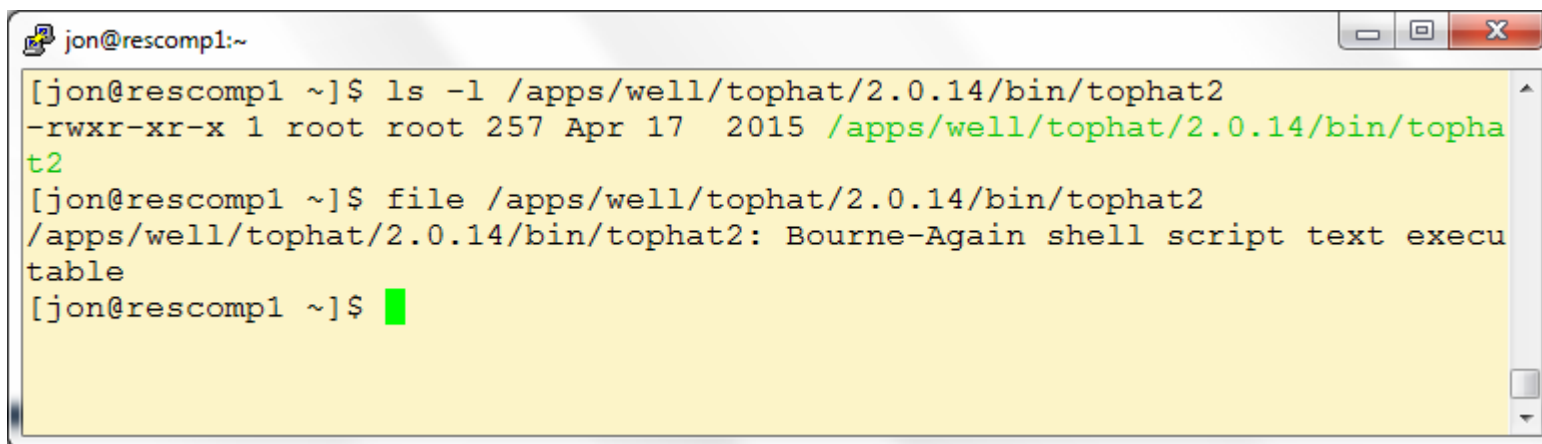
Different types of “application”

Marked by the ‘execute’ file permission bits

- different permissions for owner, group and other

Just because it is marked as executable doesn’t mean it is

- dynamic libraries tend to be marked executable
- Windows uses the executable bit to mean something else



```
jon@rescomp1:~  
[jon@rescomp1 ~]$ ls -l /apps/well/tophat/2.0.14/bin/tophat2  
-rwxr-xr-x 1 root root 257 Apr 17 2015 /apps/well/tophat/2.0.14/bin/topha  
t2  
[jon@rescomp1 ~]$ file /apps/well/tophat/2.0.14/bin/tophat2  
/apps/well/tophat/2.0.14/bin/tophat2: Bourne-Again shell script text execu  
table  
[jon@rescomp1 ~]$
```

Relevant Environment Variables

• Your ability to run an application is affected by a number of environment variables

- where your shell looks for applications
- where your applications look for libraries and packages
- how many threads your application uses
- ...

Some apply generally, some only to certain applications, some only if an application is configured in a certain way

- going to look at the most common
- by no means a complete list

PATH

Controls where your shell looks for binaries

- if a fully-specified path is not used

Colon-separated list of directory paths

- searched in turn for an executable file of the same name
- first match wins

The command ``which`` will show the full path to the first match

Some things to note:

- `'.'` (the current directory) is not on your path by default
- whilst there's essentially no limit to the length of PATH, some applications start to struggle if it gets too long

LD_LIBRARY_PATH

Controls where your dynamic binaries look for libraries

- sometimes...

Colon-separated list of directory paths

- searched in turn for a library with the same name
- first match wins

No `which` equivalent

If I ask you to set this it's because:

- of a build defect
 - possibly because I've been given a pre-built dynamic binary
- you want to try using a different but compatible library to the one the binary was built against

PYTHONPATH

- Controls where python looks for its packages

- takes priority over some but not all python system paths

Colon-separated list of directory paths

- searched in turn for a package with the same name
- first match wins

`pywhich` is the PYTHONPATH equivalent to `which`

Set this to use personally/group-wide installed packages

- e.g., to use a newer version of a package than the one in the central install

R_LIBS / R_LIBS_USER

Controls where R looks for its packages

- by default:
 - R_LIBS is unset
 - R_LIBS_USER is set to directory `'~/R/R.version$platform-library/x.y'`

Colon-separated list of directory paths

- only directories that exist when R is started are searched
- R_LIBS processed before R_LIBS_USER
- searched in turn for a package with the same name
- first match wins

Set this to use personally/group-wide installed packages

- e.g., to use a newer version of a package than the one in the central install

PERL5LIB

- Controls where perl looks for its packages

- takes priority over perl system paths

Colon-separated list of directory paths

- searched in turn for a package with the same name
- first match wins

Set this to use personally/group-wide installed packages

- e.g., to use a newer version of a package than the one in the central install

Unlike python and R, the way perl is installed means that you are likely to need to set PERL5LIB for normal use:

- `. /apps/well/perl/5.10.1/apps-well-perl-5.10.1.sh`
- `source /apps/well/perl/5.10.1/apps-well-perl-5.10.1.csh`

OTHERS

• There are lots of other environment variables that might affect your job. Some are application-specific, some are library-specific. A few I am aware of:

- **TMP**
 - A directory that can/should be used for temporary files
- **JAVA_TOOL_OPTIONS**
 - Default arguments passed to the Java JVM which can be overridden by command line arguments – see also `_JAVA_OPTIONS`
- **OMP_NUM_THREADS**
 - Number of concurrent threads that should be used by an application linked against the OpenMP parallel library
- **SWEEP_PATH**
 - Points to the top directory of the pb-jelly package install

Setting Up Your Environment

Environment variables are variables that get passed through into the environment of the application

- different from local variables, which don't get passed through

How to set an environment variable depends on which shell you use

- bash
 - `export FOO=bar`
- tcsh
 - `setenv FOO bar`

Setting Up Your Environment

- Variables can also be passed into an application's environment by putting them at the start of the line that calls the application

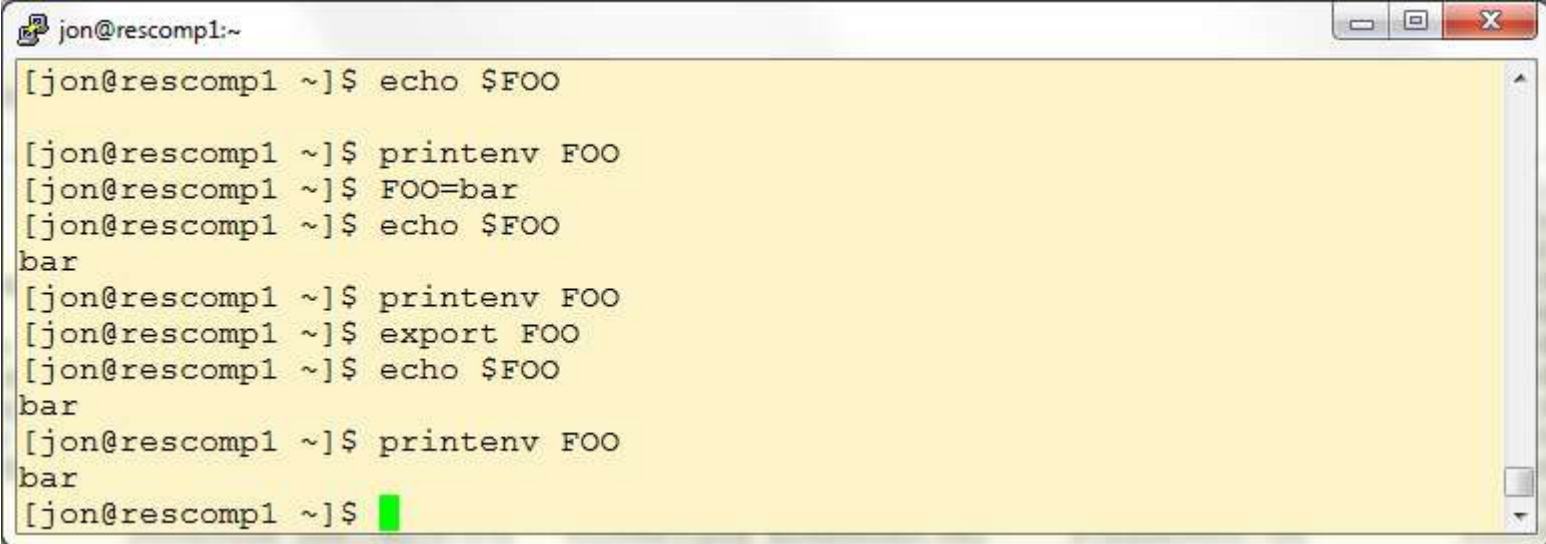
```
FOO=bar /apps/well/tophat/2.0.14/bin/tophat2
```

- note – no 'export'
- affects only that run

Setting Up Your Environment

Use ``printenv`` to check the value of an environment variable

- ``echo`` won't differentiate between local and environment variables

A terminal window titled 'jon@rescomp1:~' showing a series of commands and their outputs. The commands are: 'echo \$FOO', 'printenv FOO', 'FOO=bar', 'echo \$FOO', 'printenv FOO', 'export FOO', 'echo \$FOO', 'printenv FOO'. The outputs are: (blank), (blank), 'bar', 'bar', 'bar'. The prompt is a green cursor.

```
jon@rescomp1:~  
[jon@rescomp1 ~]$ echo $FOO  
[jon@rescomp1 ~]$ printenv FOO  
[jon@rescomp1 ~]$ FOO=bar  
[jon@rescomp1 ~]$ echo $FOO  
bar  
[jon@rescomp1 ~]$ printenv FOO  
[jon@rescomp1 ~]$ export FOO  
[jon@rescomp1 ~]$ echo $FOO  
bar  
[jon@rescomp1 ~]$ printenv FOO  
bar  
[jon@rescomp1 ~]$
```

Setting Up Your Environment

• So, back to why can't I just type, e.g., `tophat2`?

- find the version you want to use
- set your PATH environment variable appropriately

```
jon@rescomp1:~$ which tophat2
/apps/well/tophat/2.0.14/bin/tophat2
jon@rescomp1:~$ locate -b '\tophat2'
/gpfs0/apps/htseq/tophat2
/gpfs0/apps/htseq/bin/tophat2
/gpfs0/apps/htseq/head-cluster2-usr-local-genetics/bin/tophat2
/gpfs0/apps/htseq/tophat-2.0.12.Linux_x86_64/tophat2
/gpfs0/apps/well/irap/git-gcc4.7.2/bin/tophat2
/gpfs0/apps/well/irap/git-gcc4.7.2/bin/tophat2/bin/tophat2
/gpfs0/apps/well/irap/git-gcc4.7.2/tmp/tophat-2.0.10.Linux_x86_64/tophat2
/gpfs0/apps/well/tophat/2.0.11/bin/tophat2
/gpfs0/apps/well/tophat/2.0.14/bin/tophat2
jon@rescomp1:~$ export PATH=/apps/well/tophat/2.0.14/bin:$PATH
jon@rescomp1:~$ which tophat2
/apps/well/tophat/2.0.14/bin/tophat2
jon@rescomp1:~$
```

Setting Up Your Environment

• The ``module`` command simplifies setting the environment

- can set several variables at once
- can be aware of conflicts with another application, or a different version of the same application
- useful subcommands are ``avail``, ``show``, ``load`` and ``unload``

You should probably use this mechanism for

- python
- R

ShellShock bug fix has made it difficult to use with SGE

- submit with the `'-V'` option to pass environment to job
- or reactivate the module package within the script:

```
. /etc/profile.d/modules.sh
```


Setting Up Your Environment

```
jon@rescomp1:~  
[jon@rescomp1 ~]$ module avail  
  
----- /gpfs0/mgmt/modules/Modules/3.2.10/modulefiles -----  
R/2.15 intel/compiler/l_ics_2013.0.028  
R/2.15.3 intel/mpi/4.1.0.024  
R/3.0 java/1.8.0_latest  
R/3.0.3 matlab/mcr-2013b  
R/3.1 matlab/mcr-2014a  
R/3.1.0 matlab/mcr-2015a  
R/3.1.2 mpich2/3.1/gcc/4.4.7  
R/3.1.3 mvapich2/1.9/gcc/4.4.7  
R/3.1.3-atlas-3.10.2-gcc4.7.2 openmpi/1.6.5/gcc/4.4.7  
R/3.1.3-gcc4.9.3 python/2.7  
R/3.1.3-openblas-0.2.14-omp-gcc4.7.2 python/2.7.10  
R/3.2.0-openblas-0.2.14-omp-gcc4.7.2 python/2.7.10-gcc4.9.3  
R/3.2.0-openblas-0.2.14-omp-gcc4.8.2 python/2.7.10-gcc5.4.0  
R/3.2.2 python/2.7.11(default)  
R/3.2.2-openblas-0.2.14-omp-gcc4.7.2 python/2.7.6  
R/3.2.2-openblas-0.2.14-omp-gcc4.8.2 python/2.7.7-venv  
R/3.2.5 python/2.7.7-venv-unicode  
R/3.2.5-openblas-0.2.18-omp-gcc4.7.2 python/2.7.8  
R/3.2.5-openblas-0.2.18-omp-gcc4.8.2(default) python/3.4.3  
R/3.3.0 python/3.5.2-gcc5.4.0  
R/3.3.0-openblas-0.2.18-omp-gcc4.7.2 relion/1.2-gcc4.7.2  
R/3.3.0-openblas-0.2.18-omp-gcc4.8.2 relion/1.3-3-gcc4.7.2  
R/3.3.1 relion/1.3-beta3-gcc4.7.2  
R/3.3.1-openblas-0.2.18-omp-gcc4.7.2 relion/1.3-compA-gcc4.7.2  
R/3.3.1-openblas-0.2.18-omp-gcc4.8.2 relion/1.3-gcc4.7.2  
R/3.3.1-openblas-0.2.18-omp-gcc4.9.3 relion/1.3-relax-1-gcc4.7.2  
R/3.3.1-openblas-0.2.18-omp-gcc5.4.0 relion/1.3-relax-1.1-gcc4.7.2  
R/default(default) relion/1.4-fftw-3.3.4-gcc4.7.2  
amber/14-gcc4.9.1-mkl relion/1.4-gcc4.7.2  
autoconf/2.69 relion/1.4-single-fftw-3.3.4-gcc4.7.2  
cns/1.3 relion/1.4-single-gcc4.7.2
```


Setting Up Your Environment

```
jon@rescomp1:~  
[jon@rescomp1 ~]$ which python  
/usr/bin/python  
[jon@rescomp1 ~]$ python --version  
Python 2.6.6  
[jon@rescomp1 ~]$ module show python/2.7  
-----  
/gpfs0/mgmt/modules/Modules/3.2.10/modulefiles/python/2.7:  
  
module-whatis      Sets up your environment so calling python runs version 2  
.7.6  
conflict           python  
prepend-path       PATH /apps/well/python/2.7.6/bin  
prepend-path       PATH /apps/well/python/2.7/bin  
prepend-path       LD_LIBRARY_PATH /apps/well/python/2.7.6/lib  
prepend-path       LD_LIBRARY_PATH /apps/well/openblas/0.2.8-gcc4.7.2/lib  
prepend-path       LD_LIBRARY_PATH /apps/well/fftw/3.3.3-gcc4.7.2/lib  
prepend-path       PKG_CONFIG_PATH /apps/well/python/2.7.6/lib/pkgconfig  
-----  
  
[jon@rescomp1 ~]$ module load python/2.7  
[jon@rescomp1 ~]$ python --version  
Python 2.7.6  
[jon@rescomp1 ~]$
```

The Challenges Of Modular Applications

Most of the useful functionality of python, R and perl is delivered through add-on packages

- there can be multiple versions of a package but only one is going to be used
- packages tend to have dependencies on other packages
 - possibly a large number...
- packages are not necessarily compatible with one another
 - tends to really mean they have conflicting requirements for the packages/versions on which they depend

Once installed, it is hard for rescomp to update a package

- may break running code
- may break installed packages

The Challenges Of Modular Applications

Python tends to cause the most problems

- “accepted” solution is to use virtualenv, but
 - pushes the package installation problems down to users

Python package updates likely to be pushed out alongside python updates

- ``module load python/2.7.8`` has newer packages than 2.7.6, etc.
- 2.7.11 is latest available on rescomp
- will roll out 2.7.13 at some point with more updates

Diagnosing failures

Python / R / perl package missing

- try a newer version of python / R / perl
 - for python, use 2.7.10 or 2.7.11
 - module load python/2.7.11
- for R, try 3.2.5 or 3.3.1
 - module load R/3.3.1
- ask for the package to be installed
 - a link to the package's website is always useful

Diagnosing failures

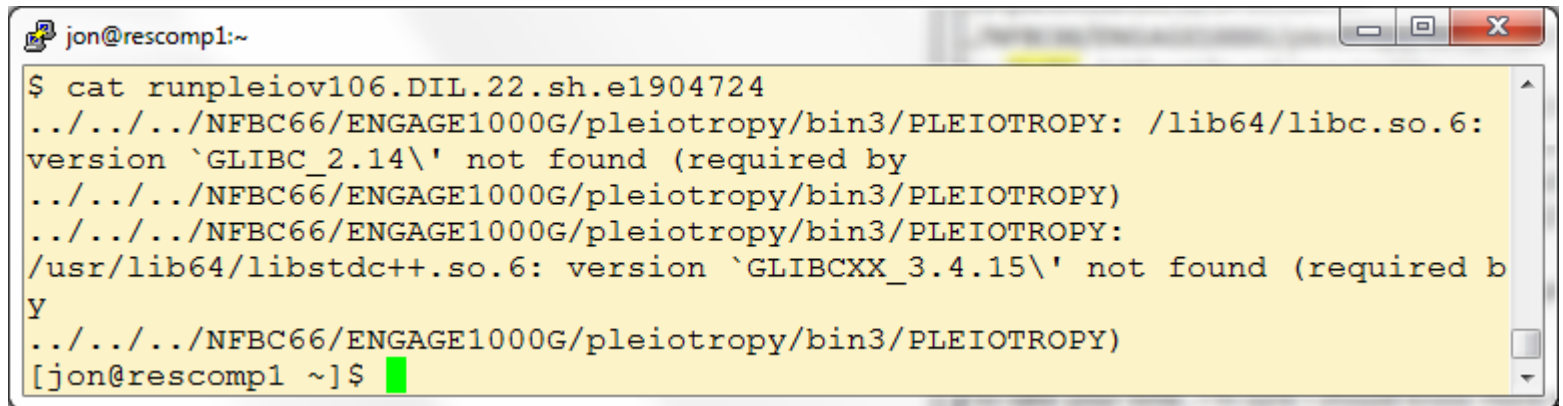
• Application not found

- use which and locate
- set your PATH environment variable appropriately
- use a fully-specified path
- ask for the application to be installed
 - a link to the application's website is always useful

Diagnosing failures

GLIBC / GLIBCXX not found

- binary was built on a box with a newer glibc/kernel than the rescomp cluster has
- rebuild required on a rescomp box
 - or a box running RHEL6 / CentOS6, etc.

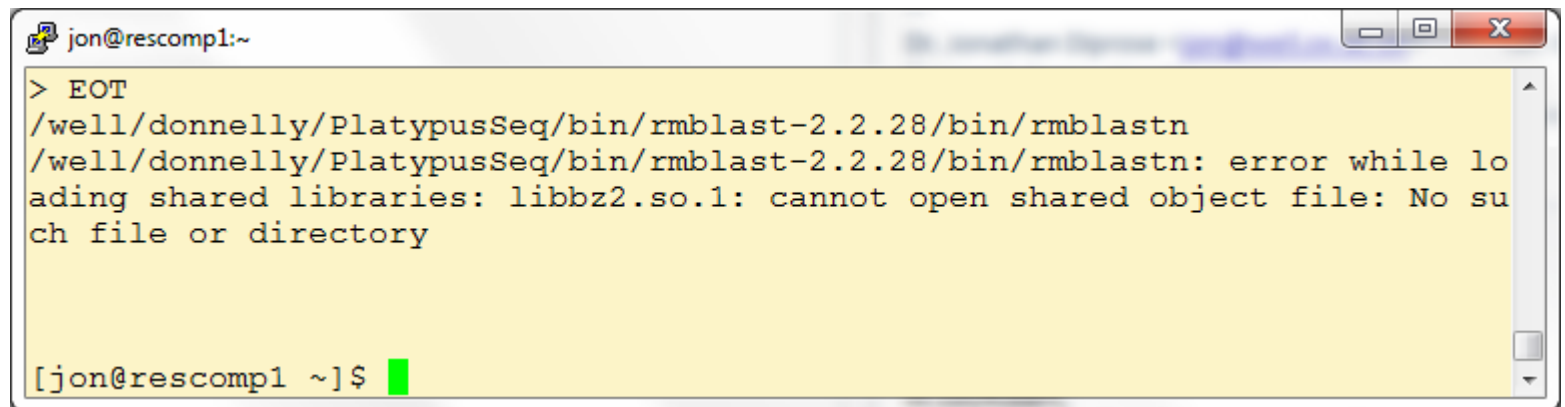


```
jon@rescomp1:~  
$ cat runpleiov106.DIL.22.sh.e1904724  
../../../../NFBC66/ENGAGE1000G/pleiotropy/bin3/PLEIOTROPY: /lib64/libc.so.6:  
version `GLIBC_2.14\' not found (required by  
../../../../NFBC66/ENGAGE1000G/pleiotropy/bin3/PLEIOTROPY)  
../../../../NFBC66/ENGAGE1000G/pleiotropy/bin3/PLEIOTROPY:  
/usr/lib64/libstdc++.so.6: version `GLIBCXX_3.4.15\' not found (required b  
Y  
../../../../NFBC66/ENGAGE1000G/pleiotropy/bin3/PLEIOTROPY)  
[jon@rescomp1 ~]$
```

Diagnosing failures

Library not found ('libSomething.so: cannot open shared object file')

- dynamic binary has failed to find a library on which it depends
- install the library if it hasn't been installed
- set LD_LIBRARY_PATH appropriately so the binary can find it
- may be necessary to rebuild

A terminal window with a yellow background and a title bar showing 'jon@rescomp1:~'. The window contains the following text:

```
> EOT
/well/donnelly/PlatypusSeq/bin/rmbblast-2.2.28/bin/rmbblastn
/well/donnelly/PlatypusSeq/bin/rmbblast-2.2.28/bin/rmbblastn: error while loading shared libraries: libbz2.so.1: cannot open shared object file: No such file or directory

[jon@rescomp1 ~]$
```